

Medical Diagnosis with Support Vector Machines

Task 1: Import Libraries

```
In [1]:
import numpy as np
import pandas as pd
from sklearn import svm
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report
```

Task 1: Get Data

```
In [2]:
column_names = ["pregnancies", "glucose", "bpressure", "skinfold", "insulin", "bmi", "pedigree", "age", "class"]
df = pd.read_csv("data.csv", names=column_names)
print(df.shape)
df.head()
```

(768, 9)

Out[2]:

	pregnancies	glucose	bpressure	skinfold	insulin	bmi	pedigree	age	class
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

Task 1: Extract Features

```
In [3]:
X = df.iloc[:, :8] #all rows all columns upto 8
X.head()
```

Out[3]:

	pregnancies	glucose	bpressure	skinfold	insulin	bmi	pedigree	age
0	6	148	72	35	0	33.6	0.627	50
1	1	85	66	29	0	26.6	0.351	31
2	8	183	64	0	0	23.3	0.672	32
3	1	89	66	23	94	28.1	0.167	21
4	0	137	40	35	168	43.1	2.288	33

Task 1: Extract Class Labels

In [4]:

```
y = df['class']
y.head()
```

Out[4]:

```
0    1
1    0
2    1
3    0
4    1
Name: class, dtype: int64
```

Task 2: Split Dataset

In [5]:

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=0)
print(X_train.shape)
print(y_train.shape)
print(X_test.shape)
print(y_test.shape)
X_test.head()
```

```
(576, 8)
(576,)
(192, 8)
(192,)
```

Out[5]:

	pregnancies	glucose	bpressure	skinfold	insulin	bmi	pedigree	age
661	1	199	76	43	0	42.9	1.394	22
122	2	107	74	30	100	33.6	0.404	23
113	4	76	62	0	0	34.0	0.391	25
14	5	166	72	19	175	25.8	0.587	51
529	0	111	65	0	0	24.6	0.660	31

Task 2: Normalize Features

In [6]:

```
scaler = StandardScaler() #normalize means we need to give a range in which it is expecting
scaler.fit(X_train)
X_train = scaler.transform(X_train)
X_train[:5,:]
```

Out[6]:

```
array([[ 1.50755225, -1.01521454, -0.40451932, -1.31921491, -0.71823753,
        -1.22070104, -0.98325882, -0.04863985],
       [-0.82986389, -0.09964691, -0.61509602,  0.9287299 ,  0.08374747,
         0.13719053, -0.62493647, -0.88246592],
       [-1.12204091, -0.95207195,  0.54307587, -1.31921491, -0.71823753,
         0.0240329 ,  0.39884168, -0.5489355 ],
       [ 2.38408331,  0.59492164,  0.64836422,  1.36583027,  2.05458297,
         0.87900167,  0.17903049,  2.03592532],
       [ 1.50755225,  0.75277813,  0.54307587,  1.55315901,  0.39089067,
         0.71555175,  0.50724171,  0.53503839]])
```

Task 3: Training a Support Vector Machine

In [7]:


```
In [10]:
```

```
clf = svm.SVC(kernel='rbf')
clf.fit(X_train, y_train)
```

```
Out[10]:
```

```
SVC()
```

Task 4: Making a single prediction

```
In [11]:
```

```
# "pregnancies", "glucose", "bpressure",
# "skinfold", "insulin", "bmi",
# "pedigree", "age", "class"
patient = np.array([[1., 50., 75., 40., 0., 45., 1.5, 20]])
patient = scaler.transform(patient)
clf.predict(patient)
```

```
Out[11]:
```

```
array([0])
```

Task 4: Testing Set Prediction

```
In [12]:
```

```
patient = np.array([X_test.iloc[0]])
patient = scaler.transform(patient) #normalizing these features
print(clf.predict(patient))
print(y_test.iloc[0])
```

```
[1]
```

```
1
```

Task 5: Accuracy on Testing Set

```
In [13]:
```

```
X_test = scaler.transform(X_test)
y_pred = clf.predict(X_test)
print(accuracy_score(y_test, y_pred))
```

```
0.7760416666666666
```

Task 5: Comparison to All-Zero Prediction

```
In [14]:
```

```
y_zero = np.zeros(y_test.shape)
print(accuracy_score(y_test, y_zero))
```

```
0.6770833333333334
```

Task 5: Precision and Recall

```
In [15]:
```

```
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.80	0.90	0.84	130
1	0.71	0.52	0.60	62

accuracy			0.78	192
macro avg	0.75	0.71	0.72	192
weighted avg	0.77	0.78	0.77	192

Class 1: PRECISION ranges from 0 to 1. If it is close to 1, our model avoids predicting that people have diabetes if they actually don't. That would be false positive. But precision may be high because we are underestimating how many people are sick.

RECALL: It also ranges from 0 to 1, if it's close to 1 our model has correctly predicted that all the people with diabetes do have disease. This is avoiding false negatives.

We need high precision and high recall

F1 SCORE: It combines the two values but focuses on the worst of the two. It's kind of pessimistic (worst) average of precision and recall

SUPPORT: It tells how many samples there were of each class.

In this project

1. I have loaded a dataset and extracted its features and labels
2. Split it into training and test subsets and normalized
3. Set up a SVM using best kernel for job and used it to make medical diagnosis

In []: